

5th Slide Set Operating Systems

Prof. Dr. Christian Baun

Frankfurt University of Applied Sciences
(1971–2014: Fachhochschule Frankfurt am Main)
Faculty of Computer Science and Engineering
christianbaun@fb2.fra-uas.de

Learning Objectives of this Slide Set

- At the end of this slide set You know/understand. . .
 - fundamental concepts of **memory management**
 - **Static partitioning**
 - **Dynamic partitioning**
 - **Buddy memory allocation**
 - how operating systems **access the memory** (address it!)
 - **Real mode**
 - **Protected mode**
 - components and concepts to implement **virtual memory**
 - Memory Management Unit (**MMU**)
 - Paged memory management (**paging**)
 - Segmented memory management (**segmentation**)
- the possible results if memory is requested
 - **Hit** and **Miss**
- the functioning and characteristics of common **replacement strategies**

Exercise sheet 5 repeats the contents of this slide set which are relevant for these learning objectives

Memory Management

- An essential function of operating systems
- Required for allocating portions of memory to programs at their request
- Also frees memory portions, which are allocated to programs, when they are not needed any longer
- 3 concepts for memory management:
 - ① **Static partitioning**
 - ② **Dynamic partitioning**
 - ③ **Buddy memory allocation**

These concepts are already somewhat older...




Image source: unknown (perhaps IBM)

A good description of the memory management concepts provides...

- **Operating Systems – Internals and Design Principles**, William Stallings, 4th edition, Prentice Hall (2001), P.305-315
- **Moderne Betriebssysteme**, Andrew S. Tanenbaum, 3rd edition, Pearson (2009), P.232-240


Concept 1: Static Partitioning

- The main memory is split into **partitions of equal size** or of **different sizes**
- Drawbacks:
 - **Internal fragmentation** occurs in any case \Rightarrow inefficient
 - The problem is moderated by partitions of different sizes, but not solved
 - The number of partitions limits the number of possible processes
 - Challenge: A process requires more memory than a partition is of size
 - Then the process must be implemented in a way that only a part of its program code is stored inside the main memory
 - When program code (modules) are loaded into the main memory *Overlay* occurs
 - \Rightarrow modules and data may become overwritten

IBM OS/360 MFT in the 1960s implemented static partitioning

Static Partitioning (1/2)

- If **partitions of equal size** are used, it does not matter which free partition is allocated to a process
 - If no partition is free, a process from main memory need to be replaced
 - The decision of which process will be replaced depends on the scheduling method (\Rightarrow slide set 8) used



Partitions
of equal size

Source:
William Stallings.
Operating Systems.
Prentice Hall. 2001

Static Partitioning (2/2)



- Processes should get a partition allocated, which fits as precise as possible
 - Objective: Little internal fragmentation
- If **partitions of different sizes** are used, 2 possible ways exist to allocate partitions to processes:

- 1 A separate process queue for each partition**

- Drawback: Some partitions may never used

- 2 A single queue for all partitions**

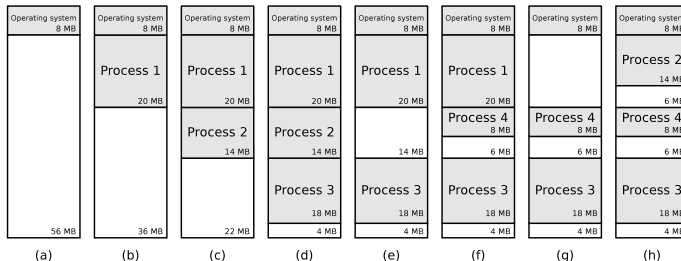
- The allocation of partitions to processes can be carried out in a flexible way
- To changed requirements of processes can be reacted quickly



Source:
William Stallings,
Operating Systems,
Prentice Hall, 2001

Concept 2: Dynamic Partitioning

- Each process gets a gapless main memory partition with the exact required size allocated



Source:
William Stallings.
Operating Systems.
Prentice Hall. 2001

- External fragmentation** occurs in any case \Rightarrow inefficient
 - Possible solution: Defragmentation
 - Requirement: Relocation of memory blocks must be supported
 - References inside processes are not allowed to become invalid because of the relocation of partitions

Implementation Concepts for Dynamic Partitioning

• First Fit

- Searches for a free block, starting from the beginning of the address space
- Quick method


• Next Fit

- Searches for a free block, starting from the latest allocation
- Fragments quickly the large area of free space at the end of the address space

• Best Fit

- Searches for the free block, which fits best
- Produces many mini-fragments and is slow

Example: A Process requires
14 MB main memory



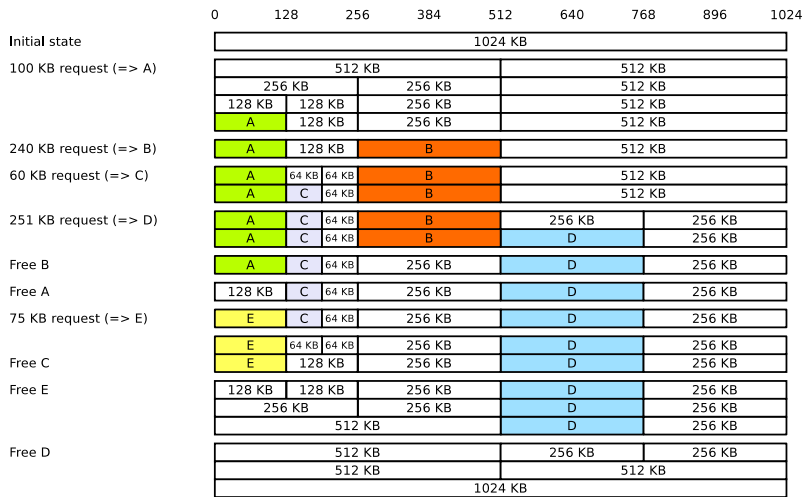
Concept 3: Buddy Memory Allocation of Donald Knuth

- Initially, a single block covers the entire memory
- If a process requires memory, the request is rounded up to the next higher power of two and an appropriate free block is searched
 - If no block of this size exists, a block of double size is searched and this block is split into 2 halves (so-called *buddies*)
 - The first block is then allocated to the process
 - If no block of double size exists, a block of four times size is searched, etc. . .
- If memory is freed, it is checked whether 2 halves of the same size can be re-combined to a larger block
 - Only previously made subdivisions are reversed!

Bddy memory management in practice

- The Linux kernel implements a variant of the buddy memory management for the page allocation
- The operating system maintains for each possible block size a list of free blocks

Buddy Memory Allocation Example



- Drawback: Internal and external fragmentation

Information about the Memory Fragmentation

- The DMA row shows the first 16 MB on a system
- The DMA32 row shows all memory > 16 MB and < 4 GB on a system
- The Normal row shows all memory > 4 GB on a system

Further information about the rows: <https://utcc.utoronto.ca/~cks/space/blog/linux/KernelMemoryZones>

```
$ cat /proc/buddyinfo
```

Node 0, zone	DMA	1	1	1	0	2	1	1	0	1	1	3
Node 0, zone	DMA32	208	124	1646	566	347	116	139	115	17	4	212
Node 0, zone	Normal	43	62	747	433	273	300	254	190	20	8	287

- column 1 \implies number of free memory chunks („buddies“) of size $2^0 * \text{PAGE_SIZE} \implies 4$ kB
- column 2 \implies number of free memory chunks („buddies“) of size $2^1 * \text{PAGE_SIZE} \implies 8$ kB
- column 3 \implies number of free memory chunks („buddies“) of size $2^2 * \text{PAGE_SIZE} \implies 16$ kB
- ...
- column 11 \implies number of free memory chunks („buddies“) of size $2^{10} * \text{PAGE_SIZE} \implies 4096$ kB = 4 MB

PAGESIZE = 4096 Bytes = 4 kB

The pagesize of a Linux system can be checked via the command: `$ getconf PAGESIZE`


Accessing Memory

- With 16-bit architectures, 2^{16} memory addresses and therefore up to 65,536 Bytes can be addressed
- With 32-bit architectures, 2^{32} memory addresses and therefore up to 4,294,967,296 Bytes = **4 GB** can be addressed
- With 64-bit architectures, 2^{64} memory addresses and therefore up to 18,446,744,073,709,551,616 Bytes = **16 Exabyte** can be addressed

!!! Question !!!

How is the memory accessed by the processes

Idea: Direct Memory Access




- Most obvious idea: Direct memory access by the processes
⇒ **Real Mode**
- Unfortunately, this is impossible in multitasking systems

Real Mode (Real Address Mode)

- Operating mode of x86-compatible CPUs
- **No memory protection**
 - Each process can access the entire memory, which can be addressed
 - Unacceptable for multitasking operating systems
- **A maximum of 1 MB main memory can be addressed**
 - Maximum main memory of an Intel 8086
 - Reason: The address bus of the 8088 contains only 20 lines
 - 20 lines \implies 20 Bits long memory addresses $\implies 2^{20} = \text{approx. 1 MB}$ memory can be addressed by the CPU
 - Only the first 640 kB (*lower memory*) can be used by the operating system (MS-DOS) and the applications
 - The remaining 384 kB (*upper memory*) contains the BIOS of the graphics card, the memory window to the graphics card memory and the BIOS ROM of the motherboard
- The term „real mode“ was introduced with the Intel 80286
 - In real mode, a CPU accesses the main memory equal to a 8086
 - Each x86-compatible CPU starts in real mode

Real Mode – Addressing

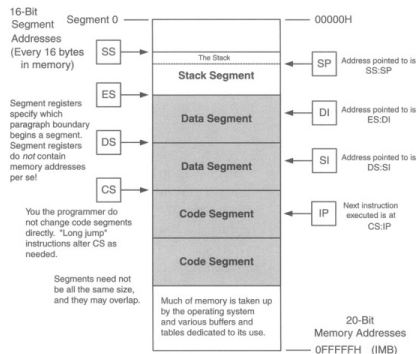
- The main memory is split into 65,536 segments
 - The memory address length is 16 Bits
 - The size of each segment is 64 Bytes ($= 2^{16} = 65,536$ bits)
- Main memory addressing is implemented via segment and offset
 - Two 16 bits long values, which are separated by a colon
Segment:Offset
 - Segment and offset are stored in the two 16-bit large registers **segment register** (= **base address register**) and **offset register** (= **index register**)
- The **segment register** stores the segments number
- The **offset register** points to an address between 0 and 2^{16} ($=65,536$), relative to the address of the segment register



Real Mode – Segment Registers since the 8086

Image Source: <http://www.c-jump.com>

- The 8086 contains 4 **segment registers**
- CS (Code Segment)
 - Segment with the source code of the program
- DS (Data Segment)
 - Segment with the global data, of the current program
- SS (Stack Segment)
 - Segment with the stack for the local data of the program
- ES (Extra Segment)
 - Segment for further data
- Since the Intel 80386, 2 further segment registers (FS, and GS) for additional extra segments exist
- The segments implement a simple form of **memory protection**



Real Mode in MS-DOS

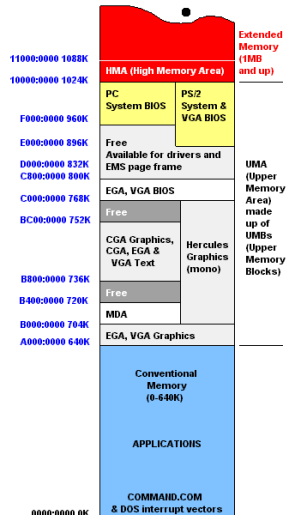
Image Source: Google Image Search

From Computer Desktop Encyclopedia
© 2001 The Computer Language Co. Inc.

```
A:\>mem
```



Memory Type	Total	Used	Free
-----	-----	-----	-----
Conventional	640K	92K	548K
Upper	0K	0K	0K
Reserved	384K	384K	0K
Extended (XMS)	742,400K	64K	742,336K
-----	-----	-----	-----
Total memory	743,424K	540K	742,884K
Total under 1 MB	640K	92K	548K
Largest executable program size		548K	(561,552 bytes)
Largest free upper memory block		0K	(0 bytes)
MS-DOS is resident in the high memory area.			

- Real mode is the default mode of MS-DOS and compatible operating systems (e.g. PC-DOS, DR-DOS and FreeDOS)



Real Mode in Microsoft Windows

- Newer operating systems only use it during the start phase and then switch to the **protected mode**



- Windows 2.0 runs only in real mode
- Windows 2.1 and 3.0 can run either in real mode or protected mode
- Windows 3.1 and later revisions run only in protected mode

Image Source: <http://virtuallyfun.superglobalmegacorp.com>

Memory Management Demands

- **Relocation**

- If processes are replaced from the main memory, it is unknown at which address they will be inserted later into the main memory again
- Finding: Processes must not refer to physical memory addresses

- **Protection**

- Memory areas must be protected against accidental or unauthorized access by other processes
- Finding: Access attempts must be verified (by the CPU)

- **Shared use**

- Despite memory protection, it must be possible for processes to collaborate via shared memory \implies slide set 10

- **Increased capacity**

- 1 MB is not enough
- It should be possible to use more memory as physically exists
- Finding: If the main memory is filled, parts of the data can be swapped

- Solution: **Protected mode** and **virtual memory**

Protected Mode


- Operating mode of x86-compatible CPUs
 - Introduced with the Intel 80286
- Increases the amount of memory, which can be addressed
 - 16-bit protected mode at 80286 \implies 16 MB main memory
 - 32-bit protected mode at 80386 \implies 4 GB main memory
- Implements the **virtual memory** concept
 - Each process is executed in its own copy of the physical address space, which is protected from other processes
 - Each process can only access its own virtual memory
 - With the Memory Management Unit (MMU), processes get address spaces provided, equal to the real mode
 - Virtual memory addresses translates the CPU with the MMU into physical memory addresses
- x86-CPU's contain 4 privilege levels (\implies slide set 7) for processes
 - Objective: Implementing memory protection to improve stability and security

Virtual Memory (1/3)

- Modern operating systems operate in protected mode
- In protected mode, the CPU supports 2 memory management methods
 - **Segmentation** exists since the 80286
 - **Paging** exists since the 80386
 - Both methods are implementation variants of the **virtual memory** concept
- Processes do not use physical memory addresses
 - This would cause issues in multitasking systems
- Instead, each process has a separate **address space**
 - This address space is an abstraction of the physical memory
 - It implements **virtual memory**
 - It is independent from the storage technology used and the existing expansion options
 - It consists of logical memory addresses, which are numbered from address 0 upwards

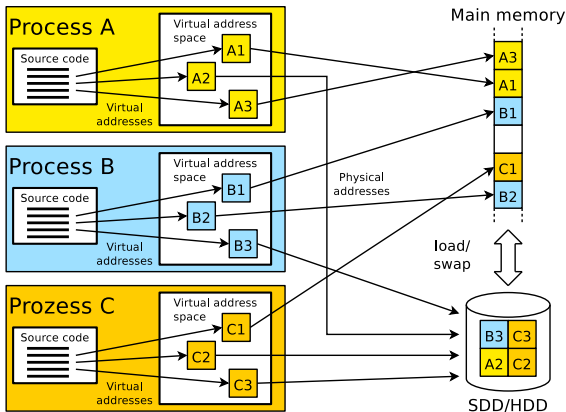
Virtual Memory (2/3)

- Address spaces can be created or erased as necessary and they are protected
 - No process can access the address space of another process without prior agreement
- The virtual memory is **mapped** to the physical memory



- With virtual memory, the main memory is utilized better
 - Processes do not need to be located in a row inside the main memory
 - Therefore, the fragmentation of the main memory is not a problem

Virtual Memory (3/3)



- Thanks to virtual memory, more memory can be accessed and used, as is physically present in the system
- **Swapping** is performed transparently for users and processes

The topic Virtual Memory is clearly explained by...

- **Betriebssysteme**, Carsten Vogt, 1st edition, Spektrum Akademischer Verlag (2001), P. 152

Paging: Paged Memory Management

- **Virtual pages** of the processes are mapped to **physical pages** in the main memory
 - All pages have the same length
 - The page size is usually 4 kb (at the Alpha architecture: 8 kB)
- Benefits:
 - No external fragmentation
 - Internal fragmentation can only occur in the last page of each process
- The operating system maintains **for each process a page table**
 - The page table stores the information where the individual pages of the process are located
- Processes only work with **virtual memory addresses**
 - Virtual memory addresses consist of 2 parts
 - The more significant part contains the page number
 - The lower significant part contains the offset (address inside a page)
 - The length of the virtual addresses is architecture dependent, and is therefore 16, 32, or 64 bits

Allocation of Process Pages to free Physical Pages

- Processes do not need to be located in a row inside the main memory
⇒ No external fragmentation


Main memory

Physical page number	0		0	A 0	0	A 0	0	A 0	0	A 0
	1		1	A 1	1	A 1	1	A 1	1	A 1
	2		2	A 2	2	A 2	2	A 2	2	A 2
	3		3	A 3	3	A 3	3	A 3	3	A 3
	4		4		4	B 0	4		4	D 0
	5		5		5	B 1	5		5	D 1
	6		6		6	B 2	6		6	D 2
	7		7		7		7	C 0	7	C 0
	8		8		8		8	C 1	8	C 1
	9		9		9		9	C 2	9	C 2
	10		10		10		10	C 3	10	C 3
	11		11		11		11		11	D 3
	12		12		12		12		12	D 4
	13		13		13		13		13	
	14		14		14		14		14	
			Load process A	Load process B	Load process C	Swap process B	Load process D			

Image source: **Operating Systems**, William Stallings, 4th edition, Prentice Hall (2001)

Address Translation by the Memory Management Unit

- Virtual memory addresses translates the CPU with the **MMU** and the page table into physical addresses
 - The operating system investigates whether the physical address belongs to the main memory or to a SSD/HDD



- If the desired data is located on the SSD/HDD, the operating system must copy the data into the main memory
- If the main memory has no more free capacity, the operating system must relocate (*swap*) data from the main memory to the SDD/HDD

The topic MMU is clearly explained by...

- **Betriebssysteme**, Carsten Vogt, 1st edition, Spektrum Akademischer Verlag (2001), P. 152-153
- **Moderne Betriebssysteme**, Andrew S. Tanenbaum, 2nd edition, Pearson (2009), P. 223-226

Implementation of the Page Table

- Impact of the page length:
 - **Short pages:** Less loss caused by internal fragmentation, but bigger page table
 - **Long pages:** Shorter page table, but more loss caused by internal fragmentation
- Page tables are stored inside the main memory

$$\text{Maximum page table size} = \frac{\text{Virtual address space}}{\text{Page size}} * \text{Size of each page table entry}$$

- Maximum page table size with 32 bit operating systems:


$$\frac{4 \text{ GB}}{4 \text{ kB}} * 4 \text{ Bytes} = \frac{2^{32} \text{ Bytes}}{2^{12} \text{ Bytes}} * 2^2 \text{ Bytes} = 2^{22} \text{ Bytes} = 4 \text{ MB}$$

- Each process in a multitasking operating system requires a page table


Page Table Structure

- Each page table record contains among others:
 - **Present bit**: Specifies whether the page is stored inside the main memory
 - **Dirty bit** (*Modified-Bit*): Specifies whether the page was modified
 - **Reference bit**: Specifies whether the page was referenced (even read operations!) \Rightarrow this is eventually relevant for the page replacement strategy used
 - **Further control bits**: Here is among others specified whether...
 - User mode processes have only read access to the page or write access too (**read/write bit**)
 - User-mode processes are allowed to access the page (**user/supervisor bit**)
 - Modifications are immediately passed down (**write-through**) or when the page is removed (**write-back**) from main memory (**write-through bit**)
 - The page may be loaded into the cache or not (**cache-disable bit**)
 - **Physical page address**: Is concatenated with the offset of the virtual address

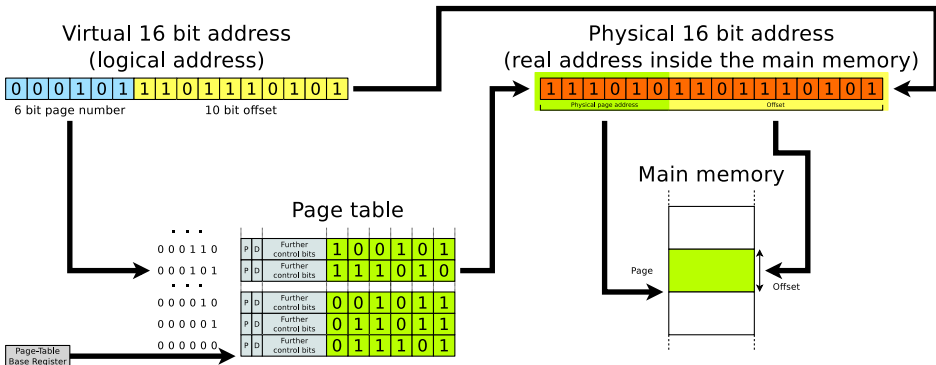
Virtual (logical) address



Page table entry



Address Translation with Paging (single level)



2 registers enable the MMU to access the page table


- **Page-Table Base Register (PTBR):** Address where the page table of the current process starts
- **Page-Table Length Register (PTLR):** Length of the page table of the current process

Address Translation with Paging (2 levels)

Virtual 32 bit address
(logical address)




Page table directory
(first layer)



Page-Table Base Register


Page table
(second layer)



Physical 32 bit address
(real address inside the main memory)



Main memory




The topic Paging is clearly explained by...

- Betriebssysteme, Eduard Glatz, 2nd edition, dpunkt (2010), P.450-457
- Betriebssysteme, William Stallings, 4th edition, Pearson (2003), S.394-399
- <http://wiki.osdev.org/Paging>

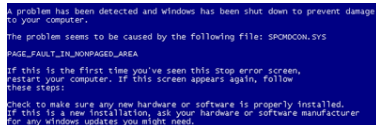
Page Fault Exception

- A program tries to access a page, which is not located in the physical main memory
 - The **present bit** in each page table record indicates whether a page is located inside main memory or not
- The operating system handles the page fault, while it carries out these steps:
 - Determine the location of the data in secondary storage (SSD/HDD)
 - Obtain free main memory pages
 - Load requested data into the pages
 - Update the page table of the process
 - Return control to the program, which retries the instruction that caused the page fault



Access Violation Exception or General Protection Fault Exception

- Also called **Segmentation fault** or **Segmentation violation**
 - A paging issue, which has nothing to do with segmentation!
- A process tries to access a virtual memory address, which it is not allowed to access \implies crash
 - Example: A process tries to carry out a write access to a read-only page



A problem has been detected and Windows has been shut down to prevent damage to your computer.

The problem seems to be caused by the following file: SPDMCON.SYS

PAGE_FAULT_IN_NONPAGED_AREA

If this is the first time you've seen this Stop error screen, restart your computer. If this screen appears again, follow these steps:

Check to make sure any new hardware or software is properly installed. If this is a new installation, ask your hardware or software manufacturer for any windows updates you might need.





Image source: Wikipedia, <http://telcontar.net/store/archive/CrashGallery/images/crash/m/crash13.png> and http://www.dtec-computers.com/images/jpg/computer_repair/blue-screen-of-death.gif

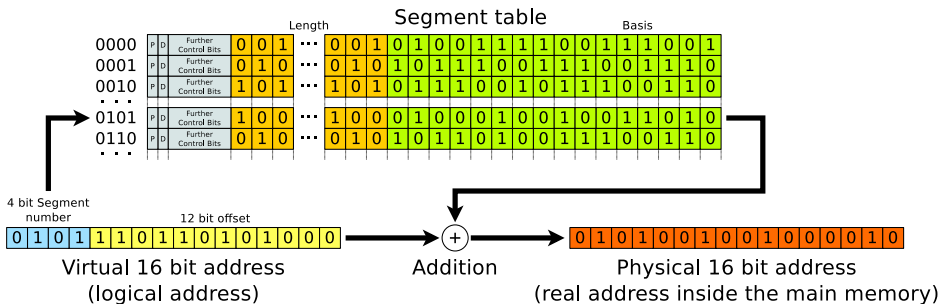
Segmentation

- A further virtual memory management method
- The virtual memory of the processes consists of **segments** of different length

The maximum segment length explains the example on the next slide

- The operating system maintains **for each process a segment table**
 - Each record in the segment table contains the length of the segment and the start address in the main memory
 - Virtual addresses of the processes are translated into physical addresses by using the segment tables
- No internal fragmentation
- External Fragmentation occurs as with dynamic partitioning
 - But not so much of it

Address Translation with Segmentation




- The **maximum segment length** is determined by the length of the offset of the virtual addresses
 - In this example, the length of the offset is 12 bits
 \Rightarrow maximum segment length = $2^{12} = 4,096$ bits = 512 Bytes

Segment Table Structure

- Each segment table record contains among others:
 - **Present bit**: Specifies whether the segment is stored inside the main memory
 - **Dirty bit** (*Accessed bit* or *Modified bit*): Specifies whether the segment was modified
 - **Further control bits**: Specifies e.g. access control privileges
 - **Length**: Length of the segment
 - **Segment basis**: Is added to the offset of the virtual address

Virtual (logical) address



Segment table entry



- If a program tries to access a segment, which is not in the physical main memory, a **segment not present** exception is raised
 - The **present bit** in each segment table record indicates whether a segment is located inside main memory or not

Summary: Real Mode and Protected Mode

• Real mode

- Operating mode of x86-compatible CPUs
- The CPU accesses the main memory equal to an Intel 8086 CPU
- No memory protection
 - Each process can access the entire main memory

• Protected mode

- Operating mode of x86-compatible CPUs
- Implements the **virtual memory** concept
- **16 bit protected mode** (introduced with the Intel 80286)
 - Only **segmentation**
 - With the 24 bits address bus, a maximum of 2^{24} Bytes = 16 MB of physical main memory can be addressed
 - Main memory is split into segments of 2^{16} Bytes = 64 kB maximum each
- **32 bit protected mode** (introduced with the Intel 80386)
 - Main memory is split into segments of 2^{32} Bytes = 4 GB maximum each
 - **Paging** can be used in addition to the segmentation
 - A maximum of 4 GB physical main memory can be addressed

Virtual Memory in Modern Operating Systems (1/2)

- Modern operating systems (for x86) operate in protected mode and use only paging
 - Segmentation is no longer used
- This method is called **flat memory model**
 - Data, code, extra and stack segment cover the entire address space
 - This way, the entire address space of each process is addressable via the offset
 - Segmentation does not provide memory protection any more
 - But this is not a problem, because of the paging
- Benefit: Operating systems can be ported more easily to other CPU architectures, which do not implement segmentation

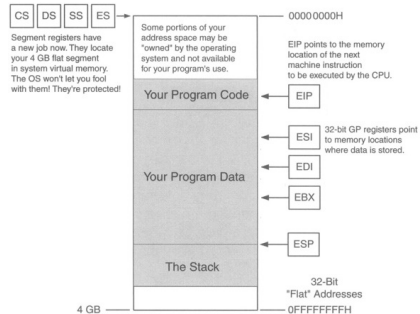


Image Source: <http://www.c-jump.com>

Virtual Memory in Modern Operating Systems (2/2)

- Some architectures:
 - IA32 (see slide 30)
 - 2-level page table
 - Length of virtual addresses: 32 bits
 - $10+10+12 \implies$ 10 bits for the 2 page tables plus 12 bits offset
 - IA32 mit Physical Address Extension (PAE) \implies Pentium Pro
 - 3-level page table
 - Length of virtual addresses: 32 bits
 - $2+9+9+12 \implies$ 2 bits for the first page table and 9 bits each for the 2 further page tables plus 12 bits offset
 - PPC64
 - 3-level page table
 - Length of virtual addresses: 41 bits
 - $10+10+9+12 \implies$ 10 bits for the first two page tables, 9 bits for the third page table plus 12 bits offset
 - AMD64 (x86-64)
 - 4-level page table
 - Length of virtual addresses: 48 bits
 - $9+9+9+9+12 \implies$ 9 bits each for the 4 page tables plus 12 bits offset

Hit Rate and Miss Rate

- In case of a request to a computer memory, 2 results are possible:
 - **Hit**: Requested data is available
 - **Miss**: Requested data is missing
- 2 Key figures are used to evaluate the efficiency of a computer memory
 - **Hit rate**: The number of requests to the computer memory, with result in hit, divided by the total number of requests
 - Result is between 0 and 1
 - The greater the value, the better is the efficiency of the computer memory
 - **Miss rate**: The number of requests to the computer memory, with result in miss, divided by the total number of requests
 - $\text{Miss rate} = 1 - \text{hit rate}$

Page Replacement Strategies

- It makes sense to keep the data (\implies **pages**) inside main memory, which is frequently accessed
- Some **replacement strategies**:
 - **OPT** (Optimal strategy)
 - **LRU** (Least Recently Used)
 - **LFU** (Least Frequently Used)
 - **FIFO** (First In First Out)
 - **Clock / Second Chance**
 - **TTL** (Time To Live)
 - **Random**


A well understandable explanation of the page replacement strategies. . .

- OPT, FIFO, LRU and Clock provides **Operating Systems**, *William Stallings*, 4th edition, Prentice Hall (2001), P.355-363
- FIFO, LRU, LFU and Clock provides **Betriebssysteme**, *Carsten Vogt*, 1st edition, Spektrum Verlag (2001), P.162-163
- FIFO, LRU and Clock provides **Moderne Betriebssysteme**, *Andrew S. Tanenbaum*, 2nd edition, Pearson (2009), P.237-242
- FIFO, LRU, LFU and Clock provides **Betriebssysteme**, *Eduard Glatz*, 2nd edition, dpunkt (2010), P.471-476

Optimal strategy (OPT)

Image Source: Lukasfilm Games

- Replaces the page, which is **not accessed for the longest time in the future**
- Impossible to implement!
 - Reason: Nobody can predict the future
 - Therefore, the operating system must take into account the past
- OPT is used to evaluate the efficiency of other replacement strategies



Requests: **1 2 3 4 1 2 5 1 2 3 4 5**

Page 1:	1	1	1	1	1	1	1	1	3	3	3
Page 2:		2	2	2	2	2	2	2	2	4	4
Page 3:			3	4	4	4	5	5	5	5	5

→ 7 Miss

The **requests** are requests for pages inside the virtual address space of a process. If the requested page is not inside the cache, it is read from the main memory or the swap

Least Recently Used (LRU)

- Replaces the page, which was **not accessed for the longest time**
- All pages are referenced in a queue
- If a page is loaded into memory or referenced, it is moved to the front of the queue
- If the memory has no more free capacity and a miss occurs, the page at the end of the queue is replaced
- Drawback: Ignores the number of accesses

Requests: **1 2 3 4 1 2 5 1 2 3 4 5**

Page 1:	1	1	1	2	3	4	1	2	5	1	2	3
Page 2:		2	2	3	4	1	2	5	1	2	3	4
Page 3:			3	4	1	2	5	1	2	3	4	5

→ 10 Miss

Least Frequently Used (LFU)

- Replaces the page, which was **least often accessed**
- For each page in memory, a reference counter exists in the page table, which stores the number of accesses
- If the memory has no more free capacity and a miss occurs, the page is replaced, which has the lowest value in its reference counter
- Benefit: Takes into account the number of accesses
- Drawback: Pages which have been accessed often in the past, may block the memory

Requests: **1 2 3 4 1 2 5 1 2 3 4 5**

Page 1:	<div><div>1</div><div>1</div><div>1</div><div>4</div><div>4</div><div>4</div><div>5</div><div>5</div><div>5</div><div>3</div><div>4</div><div>5</div></div>
Page 2:	<div><div></div><div>2</div><div>2</div><div>2</div><div>1</div><div>1</div><div>1</div><div>1</div><div>1</div><div>1</div><div>1</div><div>1</div></div>
Page 3:	<div><div></div><div></div><div>3</div><div>3</div><div>3</div><div>2</div><div>2</div><div>2</div><div>2</div><div>2</div><div>2</div><div>2</div></div>

→ 10 Miss

First In First Out (FIFO)

- Replaces the page, which is stored **in memory for the longest time**
- Assumption: increasing the memory results in fewer or, at worst, the same miss number
- Problem: Laszlo Belady demonstrated in 1969 that for certain access patterns, FIFO creates with an increased memory capacity more miss (⇒ **Belady's anomaly**)
 - Until the discovery of Belady's Anomaly, FIFO was considered a good replacement strategy

Belady's Anomaly (1969)

Requests: **1 2 3 4 1 2 5 1 2 3 4 5**

Page 1:	1	1	1	4	4	4	5	5	5	5	5
Page 2:		2	2	2	1	1	1	1	1	3	3
Page 3:			3	3	3	2	2	2	2	2	4

→ 9 Miss

Page 1:	1	1	1	1	1	1	5	5	5	5	4
Page 2:		2	2	2	2	2	2	1	1	1	1
Page 3:			3	3	3	3	3	3	2	2	2
Page 4:				4	4	4	4	4	4	3	3

→ 10 Miss

More information about Belady's anomaly

Belady, Nelson and Shedler. *An Anomaly in Space-time Characteristics of Certain Programs Running in a Paging Machine*. Communications of the ACM. Volume 12 Issue 6. June 1969

Clock / Second Chance

- This strategy uses the *reference bit* (see slide 28), which exists in the page table for each page
 - If a page is loaded into memory \Rightarrow reference bit = 0
 - If a page is accessed \Rightarrow reference bit = 1
- A pointer indicates the last accessed page
- In case of a miss, the memory is searched from the position of the pointer for the first page, whose reference bit has value 0
 - This page will be replaced
 - For all pages, which are examined during the searching, where the reference bit has value 1, it is set to value 0

Requests: **1 2 3 4 1 2 5 1 2 3 4 5**

Page 1:	1^x 0	1 0	1 0	4^x 0	4 0	4 0	5^x 0	5 0	5 0	3^x 0	4^x 0	4 0
Page 2:		2^x 0	2 0	2 0	1^x 0	1 0	1 1	1^x 1	1 1	1 1	1 0	5^x 0
Page 3:			3^x 0	3 0	3 0	2^x 0	2 0	2 2	2^x 1	2 2	2 0	2 0

→ 10 Miss

Further Replacement Strategies

- **TTL** (Time To Live): Each page gets a time to live value, when it is stored in the memory
 - If the TTL has exceeded, the page can be replaced
- **Random**: Random pages are replaced
 - Benefits Simple and resource-saving replacement strategy
 - Reason: No need to store information about the access behavior

The random replacement strategy is (was) used in practice

- The operating systems IBM OS/390 and Windows NT 4.0 on SMP systems use the random replacement strategy
- The Intel i860 RISC CPU uses the Random replacement strategy for the cache