

7th Slide Set

Operating Systems

Prof. Dr. Christian Baun

Frankfurt University of Applied Sciences
(1971–2014: Fachhochschule Frankfurt am Main)
Faculty of Computer Science and Engineering
christianbaun@fb2.fra-uas.de

Learning Objectives of this Slide Set

- At the end of this slide set You know/understand...
 - what a **process** is from operating system perspective
 - what information the **process context** contains in detail
 - **User context**
 - **Hardware context**
 - **System context**
 - the different process states by discussing **process state models**
 - how **process management works** in detail with **process tables**, **process control blocks** and **status lists**
 - what steps operating systems perform when **processes are created** (via fork or exec) or **erased**
 - what **fork bombs** are
 - the **structure of UNIX processes in memory**
 - what **System calls** are and how they work

Exercise sheet 7 repeats the contents of this slide set which are relevant for these learning objectives

Process and Process Context

We already know...

- A **process** (lat. *procedere* = proceed, move forward) is an instance of a program that is running
 - Processes are dynamic objects and represent sequential activities in a computer system
 - On computers, all the time, multiple processes are executed
 - In multitasking mode, the CPU is switched back and forth between the processes
-
- A process includes in addition to the program code its **context**
 - 3 types of context information manages the operating system:
 - **User context**
 - Content of the allocated address space (virtual memory) \implies slide set 5
 - **Hardware context** (\implies slide 4)
 - CPU registers
 - **System context** (\implies slide 5)
 - Information, which stores the operating system about a process
 - The operating system stores the information of the hardware context and system context in the **process control block** (\implies slide 6)

Hardware Context

- The **hardware context** is the content of the CPU registers during process execution
- Registers whose content needs to be backed up in the event of a process switch:
 - Program Counter (Instruction Pointer) – stores the memory address of the next instruction to be executed
 - Stack pointer – stores the address at the current end of the stack
 - Base pointer – points to an address in the stack
 - Instruction register – stores the instruction, which is currently executed
 - Accumulator – stores operands for the ALU and their results
 - Page-table base Register – stores the address of the page table of the running process
 - Page-table length register – stores the length of the page table of the running process

Some of these registers have been discussed in slide set 3 and slide set 5


System Context

- The **system context** is the information, the operating system stores about a process
- Examples:
 - Record in the process table
 - Process ID (PID)
 - Process state
 - Information about parent or child processes
 - Priority
 - Identifiers - access credentials to resources
 - Quotas - allowed usage quantity of individual resources
 - Runtime
 - Opened files
 - Assigned devices


Process Table and Process Control Blocks

- Each process has its own process context, which is independent of the contexts of other processes
- For managing the processes, the operating system implements the **process table**
 - It is a list of all existing processes.
 - It contains for each process a record which is called **process control block**

Process table

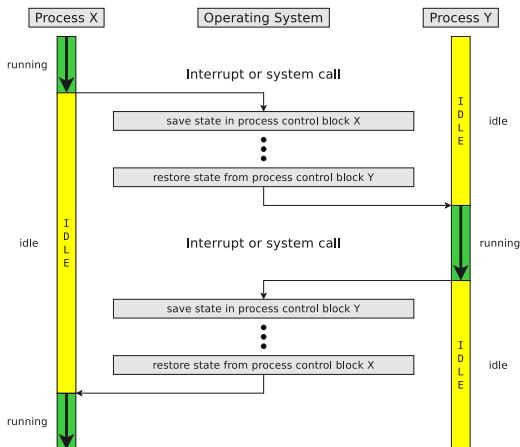


Process control blocks



Process Switching

- If the CPU is switched from one process to another one, the context (\Rightarrow CPU register content) of the running process is stored in the process control block
 - If a process gets the CPU assigned, its context gets restored by using the content of the process control block



- Each process is at any moment in a particular **state**
 \Rightarrow Process state models

Process States

We already know...

Every process is at any moment in a state

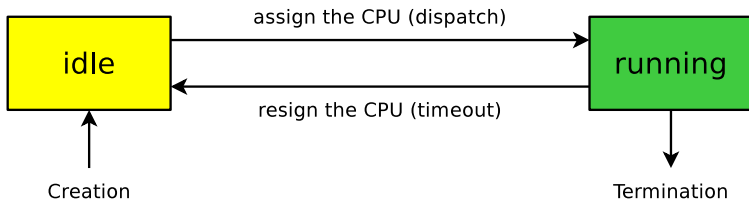
- The number of different states depends on the process state model of the operating system used

Question

How many process states must a process model contain at least?


Process State Model with 2 States

- In principle, 2 process states are enough
 - **running**: The CPU is allocated to a process
 - **idle**: The processes waits for the allocation of CPU



Process State Model with 2 States (Implementation)

- The processes in **idle** state must be stored in a queue, in which they wait for execution
 - The list is sorted according to the process priority or waiting time



The priority (proportional computing power) in Linux has a value from -20 to +19 (in integer steps). -20 is the highest priority and 19 is the lowest priority. The default priority is 0. Normal users can assign priorities from 0 to 19. The system administrator (root) can assign negative values too.

- This model also shows the working method of the **dispatcher**
 - The job of the dispatcher is to carry out the state transitions
- The execution order of the processes is specified by the **scheduler**, which uses a **scheduling algorithm** (see slide set 8)

Conceptual Error of the Process State Model with 2 States

- The process state model with 2 states assumes that all processes are ready to run at any time
 - This is unrealistic!
 - Almost always do processes exist, which are **blocked**
 - Possible reasons:
 - They wait for the input or output of an I/O device
 - They wait for the result of another process
 - They wait for a user reaction (interaction)
 - Solution: The idle processes be categorized into 2 groups
 - Processes, which are **ready**
 - Processes, which are **blocked**
- ⇒ Process state model with 3 states

Process State Model with 3 States

- Each process is in one of the following states:

- running:**


- The CPU is assigned to the process and executes its instructions

- ready:**

- The process could immediately execute its instructions on the CPU and it is currently waiting for the allocation of the CPU

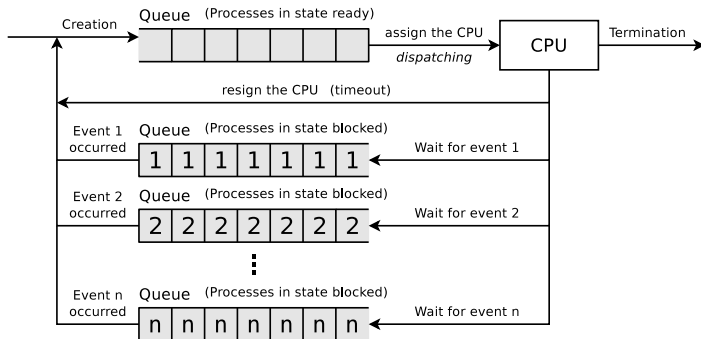
- blocked:**

- The process can currently not be executed and is waiting for the occurrence of an event or the satisfaction of a condition
- This may be e.g. a message of another process or of an input/output device or the occurrence of a synchronization event



Process State Model with 3 States – Implementation


- In practice, operating systems (e.g. Linux) implement multiple queues for processes in **blocked** state



- During state transition, the process control block of the affected process is removed from the old status list and inserted into the new status list
- No separate list exists for processes in **running** state

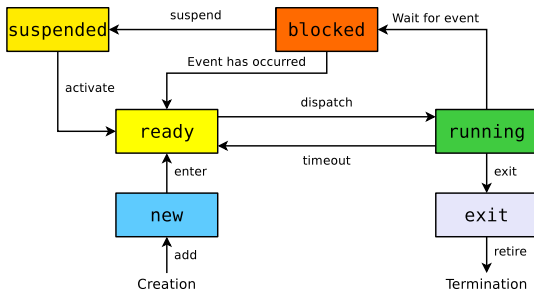
Process State Model with 5 States

- It makes sense to expand the process state model with 3 states by 2 further process states
 - new**: The process (process control block) has been created by the operating system but the process is not yet added to the queue of processes in **ready** state
 - exit**: The execution of the process has finished or was terminated, but for various reasons the process control block still exists
- Reason for the existence of the process states **new** and **exit**:
 - On some systems, the number of executable processes is limited in order to save memory and to specify the degree of multitasking



Process State Model with 6 States

- If not enough physical main memory capacity exists for all processes, parts of processes must be swapped out \Rightarrow **swapping**
- The operating system outsources processes, which are in **blocked** state
- As a result, more main memory capacity is available for the processes in the states **running** and **ready**
 - Therefore it makes sense to extend the process state model with 5 states with a further process state **suspended**



Process Creation in Linux/UNIX via `fork` (1/2)

- The system call `fork()` is the only way to create a new process
- If a process calls `fork()`, an identical copy is started as a new process
 - The calling process is called **parent process**
 - The new process is called **child process**
- The child process has after creation the same source code
 - Also the program counters have the same value, which means they refer to the same source code line
- Opened files and memory areas of the parent process are copied for the child process and are independent from the parent process
 - Child process and parent process both have their own process context

`vfork` is a variant of `fork`, which does not copy the address space of the parent process, and therefore causes less overhead than `fork`. Using `vfork` is useful if the child process is to be replaced by another process immediately after its creation. In this course `vfork` is not further discussed.

Process Creation in Linux/UNIX via fork (2/2)

- If a process calls `fork()`, an exact copy is created
 - The processes differ only in the return values of `fork()`

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <stdlib.h>
4
5 void main() {
6     int return_value = fork();
7
8     if (return_value < 0) {
9         // If fork() returns -1, an error happened.
10        // Memory or processes table have no more free capacity.
11        ...
12    }
13    if (return_value > 0) {
14        // If fork() returns a positive number, we are in the parent process.
15        // The return value is the PID of the newly created child process.
16        ...
17    }
18    if (return_value == 0) {
19        // If fork() returns 0, we are in the child process.
20        ...
21    }
22 }
```

Process Tree

- By creating more and more new child processes with `fork()`, a tree of processes (\Rightarrow process hierarchy) is created
- The command `ps tree` returns an overview about the processes, running in Linux/UNIX, as a tree according to their parent/child relationships

```
$ ps tree
init--Xprt
  |-acpid
  ...
  |-gnome-terminal--4*[bash]
    |
    |   |-bash---su---bash
    |   |-bash--gv---gs
    |   |
    |   |   |-pstree
    |   |   |
    |   |   |   |-xterm---bash---xterm---bash
    |   |   |   |-xterm---bash---xterm---bash---xterm---bash
    |   |   |   `--xterm---bash
    |   |   |-gnome-pty-helpe
    |   |   `--{gnome-terminal}
    |-4*[gv---gs]
```

Information about processes in Linux/UNIX

```
$ ps -eFw
```

| UID | PID | PPID | C | SZ | RSS | PSR | STIME | TTY | TIME | CMD |
|------|-------|------|---|--------|--------|-----|-------------|-----|----------|---------------------------------|
| root | 1 | 0 | 0 | 51286 | 7432 | 2 | Apr11 ? | | 00:00:03 | /sbin/init |
| root | 1073 | 1 | 0 | 90930 | 6508 | 0 | Apr11 ? | | 00:00:00 | /usr/sbin/lightdm |
| root | 1551 | 1073 | 0 | 60913 | 6772 | 2 | Apr11 ? | | 00:00:00 | lightdm --session-child 14 23 |
| bnc | 2143 | 1551 | 0 | 1069 | 1560 | 0 | Apr11 ? | | 00:00:00 | /bin/sh /etc/xdg/xfce4/xinitrc |
| bnc | 2235 | 2143 | 0 | 85195 | 18888 | 3 | Apr11 ? | | 00:00:11 | xfce4-session |
| bnc | 2284 | 2235 | 0 | 110875 | 45256 | 3 | Apr11 ? | | 00:06:20 | xfce4-panel --display :0.0 |
| bnc | 2389 | 2235 | 0 | 129173 | 47904 | 0 | Apr11 ? | | 00:00:26 | xfce4-terminal --geometry=80x24 |
| bnc | 2471 | 2389 | 0 | 5374 | 5360 | 2 | Apr11 pts/0 | | 00:00:00 | bash |
| bnc | 2487 | 1 | 5 | 316370 | 395892 | 0 | Apr14 ? | | 00:08:58 | /opt/google/chrome/chrome |
| bnc | 2525 | 2389 | 0 | 5895 | 6620 | 3 | Apr11 pts/5 | | 00:00:00 | bash |
| bnc | 3105 | 2284 | 0 | 597319 | 257520 | 0 | Apr11 ? | | 00:05:22 | kate -b |
| bnc | 3122 | 3105 | 0 | 5364 | 5156 | 2 | Apr11 pts/6 | | 00:00:00 | /bin/bash |
| bnc | 11196 | 2471 | 0 | 269491 | 181048 | 0 | Apr14 pts/0 | | 00:00:25 | okular bsrn_vorlesung_04.pdf |
| bnc | 16325 | 1 | 0 | 346638 | 146872 | 3 | 10:31 ? | | 00:00:16 | evince BA.pdf |
| bnc | 17384 | 2525 | 1 | 223478 | 61312 | 2 | 10:39 pts/5 | | 00:00:49 | dia |
| bnc | 19561 | 2471 | 0 | 9576 | 3340 | 0 | 11:20 pts/0 | | 00:00:00 | ps -eFw |

- C (CPU) = CPU utilization of the process in percent
- SZ (Size) = virtual process size = Text segment, heap and stack (see slide 31)
- RSS (Resident Set Size) = Occupied physical memory (without swap) in kB
- PSR = CPU core assigned to the process
- STIME = start time of the process
- TTY (Teletypewriter) = control terminal. Usually a virtual device: pts (pseudo terminal slave)
- TIME = consumed CPU time of the process (HH:MM:SS)

Independent of Parent and Child Processes

- The example demonstrates that parent and child processes operate independently of each other and have different memory areas

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <stdlib.h>
4
5 void main() {
6     int i;
7     if (fork())
8         // Parent process source code
9         for (i = 0; i < 5000000; i++)
10             printf("\n Parent: %i", i);
11     else
12         // Child process source code
13         for (i = 0; i < 5000000; i++)
14             printf("\n Child : %i", i);
15 }
```

```
Child : 0
Child : 1
...
Child : 21019
Parent: 0
...
Parent: 50148
Child : 21020
...
Child : 129645
Parent: 50149
...
Parent: 855006
Child : 129646
...
```

- The output demonstrates the CPU switches between the processes
- The value of the loop variable `i` proves that parent and child processes are independent of each other
 - The result of execution can not be reproduced

Execute on a single CPU core only...

```
$ taskset --cpu-list 1 ./fork_beispiel2.c
```

The PID Numbers of Parent and Child Process (1/2)

```
1 #include <stdio.h>
2 #include <unistd.h>
3 #include <stdlib.h>
4
5 void main() {
6     int pid_of_child;
7
8     pid_of_child = fork();
9
10    // An error occurred --> program abort
11    if (pid_of_child < 0) {
12        perror("\n fork() caused an error!");
13        exit(1);
14    }
15
16    // Parent process
17    if (pid_of_child > 0) {
18        printf("\n Parent: PID: %i", getpid());
19        printf("\n Parent: PPID: %i", getppid());
20    }
21
22    // Child process
23    if (pid_of_child == 0) {
24        printf("\n Child:  PID: %i", getpid());
25        printf("\n Child:  PPID: %i", getppid());
26    }
27 }
```

- This example creates a child process
- Child process and parent process both print:
 - Own PID
 - PID of parent process (PPID)

The PID Numbers of Parent and Child Process (2/2)

- The output is usually similar to this one:

```
Parent: PID: 20835
Parent: PPID: 3904
Child:  PID: 20836
Child:  PPID: 20835
```

- This result can be observed sometimes:

```
Parent: PID: 20837
Parent: PPID: 3904
Child:  PID: 20838
Child:  PPID: 1
```


- The parent process was terminated before the child process
 - If a parent process terminates before the child process, it gets `init` as the new parent process assigned
 - Orphaned processes are always adopted by `init`

`init` (PID 1) is the first process in Linux/UNIX

All running processes originate from `init` \implies `init` = father of all processes

Replacing Processes via exec

- The system call `exec()` replaces a process with another one
 - A **concatenation** takes place
 - The new process gets the PID of the calling process
- If the objective is to start a new process out a program, it is necessary, to create a new process with `fork()` and to replace this new process with `exec()`
 - If no new process is created with `fork()` before `exec()` is called, the parent process gets lost
- Steps of a program execution from a shell:
 - The shell creates with `fork()` an identical copy of itself
 - In the new process, the actual program is started with `exec()`



exec Example

```
$ ps -f
UID          PID    PPID    C  STIME TTY          TIME CMD
user         1772    1727     0  May18 pts/2        00:00:00 bash
user        12750    1772     0  11:26 pts/2        00:00:00 ps -f
$ bash
$ ps -f
UID          PID    PPID    C  STIME TTY          TIME CMD
user         1772    1727     0  May18 pts/2        00:00:00 bash
user        12751    1772    12  11:26 pts/2        00:00:00 bash
user        12769   12751     0  11:26 pts/2        00:00:00 ps -f
$ exec ps -f
UID          PID    PPID    C  STIME TTY          TIME CMD
user         1772    1727     0  May18 pts/2        00:00:00 bash
user        12751    1772     4  11:26 pts/2        00:00:00 ps -f
$ ps -f
UID          PID    PPID    C  STIME TTY          TIME CMD
user         1772    1727     0  May18 pts/2        00:00:00 bash
user        12770    1772     0  11:27 pts/2        00:00:00 ps -f
```

- Because of the exec, the ps -f command replaced the bash and got its PID (12751) and PPID (1772)

A further exec Example

```
1 #include <stdio.h>
2 #include <unistd.h>
3
4 int main() {
5     int pid;
6     pid = fork();
7
8     // If PID!=0 --> Parent process
9     if (pid) {
10        printf("...Parent process...\n");
11        printf("[Parent] Own PID:          %d\n", getpid());
12        printf("[Parent] PID of the child:  %d\n", pid);
13    }
14    // If PID=0 --> Child process
15    else {
16        printf("...Child process...\n");
17        printf("[Child] Own PID:          %d\n", getpid());
18        printf("[Child] PID of the parent: %d\n", getppid());
19
20        // Current program is replaced by "date"
21        // "date" will be the process name in the process table
22        execl("/bin/date", "date", "-u", NULL);
23    }
24    printf("[%d ]Program abort\n", getpid());
25    return 0;
26 }
```

- The system call `exec()` does not exist as wrapper function
- But multiple variants of the `exec()` function exist
- One of these variants is `execl()`

Helpful overview about the different variants of the `exec()` function

<http://www.cs.uregina.ca/Links/class-info/330/Fork/fork.html>

Explanation of the exec Example

```
$ ./exec_example
...Parent process...
[Parent] Own PID:          25646
[Parent] PID of the child:  25647
[25646] Program abort
...Child process...
[Child] Own PID:           25647
[Child] PID of the parent: 25646
Di 24. Mai 17:25:31 CEST 2016
```

```
$ ./exec_example
...Parent process...
[Parent] Own PID:          25660
[Parent] PID of the child:  25661
[25660] Program abort
...Child process...
[Child] Own PID:           25661
[Child] PID of the parent:  1
Di 24. Mai 17:26:12 CEST 2016
```


- After printing its PID via `getpid()` and the PID of its parent process via `getppid()`, the child process is replaced via `date`
- If the parent process of a process terminates before the child process, the child process gets `init` as new parent process assigned

Since Linux Kernel 3.4 (2012) and Dragonfly BSD 4.2 (2015), it is also possible that other processes than `PID=1` become the new parent process of an orphaned process
<http://unix.stackexchange.com/questions/149319/new-parent-process-when-the-parent-process-dies/177361#177361>


3 possible Ways to create a new Process

- **Process forking:** A running process creates with `fork()` a new, identical process
- **Process chaining:** A running process creates with `exec()` a new process and terminates itself this way because it gets replaced by the new process
- **Process creation:** A running process creates with `fork()` a new, identical process, which replaces itself via `exec()` by a new process


Process forking



Process chaining



Process creation



Have Fun with Fork Bombs

- A fork bomb is a program, which calls the `fork` system call in an infinite loop
- Objective: Create copies of the process until there is no more free memory
 - The system becomes unusable

Python fork bomb

```
1 import os
2
3 while True:
4     os.fork()
```

C fork bomb

```
1 #include <unistd.h>
2
3 int main(void)
4 {
5     while(1)
6         fork();
7 }
```

PHP fork bomb

```
1 <?php
2 while(true)
3     pcntl_fork();
4 ?>
```

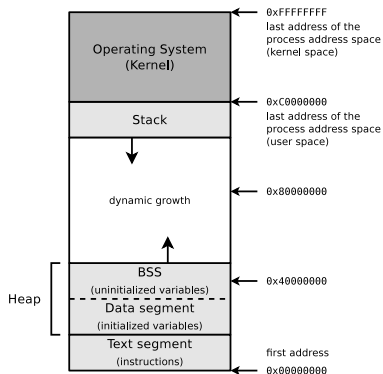
- Only protection option: Limit the maximum number of processes and the maximum memory usage per user

Structure of a UNIX Process in Memory (1/6)

- Default allocation of the virtual memory on a Linux system with a 32-bit CPU
 - 1 GB for the system (kernel)
 - 3 GB for the running process

The structure of processes on 64 bit systems is not different from 32 bit systems. Only the address space is larger and thus the possible extension of the processes in the memory.

- The **text segment** contains the program code (machine code)
- Can be shared by multiple processes
 - Must be stored for this reason only once in physical memory
 - Is therefore usually read-only
- `exec()` reads the text segment from the program file

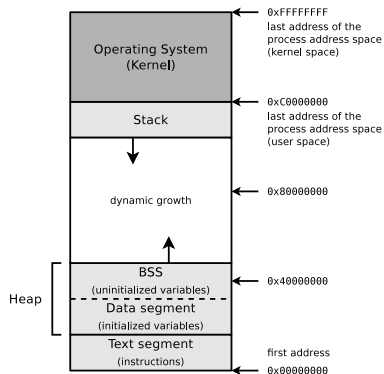


Sources

UNIX-Systemprogrammierung, *Helmut Herold*, Addison-Wesley (1996), P.345-347
Betriebssysteme, *Carsten Vogt*, Spektrum (2001), P.58-60
Moderne Betriebssysteme, *Andrew S. Tanenbaum*, Pearson (2009), P.874-877

Structure of a UNIX Process in Memory (2/6)

- The **heap** grows dynamically and consists of 2 parts:
 - 1 **data segment**
 - 2 **BSS**
- The **data segment** contains **initialized** variables and constants
 - Contains all data, which get values assigned in global declarations (outside of functions)
 - Example: `int sum = 0;`
 - `exec()` reads the data segment from the program file

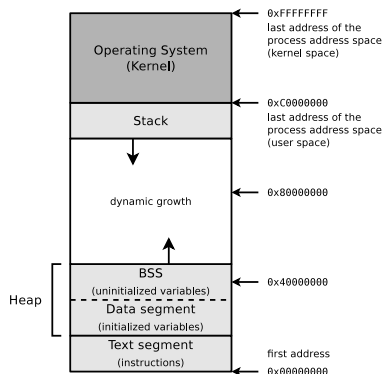


Sources

UNIX-Systemprogrammierung, *Helmut Herold*, Addison-Wesley (1996), P.345-347
Betriebssysteme, *Carsten Vogt*, Spektrum (2001), P.58-60
Moderne Betriebssysteme, *Andrew S. Tanenbaum*, Pearson (2009), P.874-877

Structure of a UNIX Process in Memory (3/6)

- The area **BSS** (*block started by symbol*) contains **uninitialized** variables
- Contains global variables (declaration is outside of functions), which get no initial values assigned
 - Example: `int i;`
- Moreover, the process can dynamically allocate memory in this area at runtime
 - In C with the function `malloc()`
- `exec()` initializes all variables in the BSS with 0




Sources

UNIX-Systemprogrammierung, *Helmut Herold*, Addison-Wesley (1996), P.345-347
Betriebssysteme, *Carsten Vogt*, Spektrum (2001), P.58-60
Moderne Betriebssysteme, *Andrew S. Tanenbaum*, Pearson (2009), P.874-877

Structure of a UNIX Process in Memory (4/6)

- The **stack** is used to implement nested function calls
 - It also contains command line arguments of the program call and environment variables
- Operates according to the LIFO (Last In First Out) principle

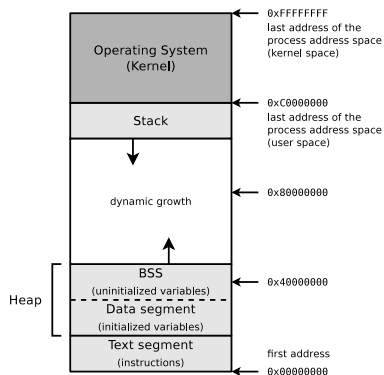


Sources

UNIX-Systemprogrammierung, *Helmut Herold*, Addison-Wesley (1996), P.345-347
Betriebssysteme, *Carsten Vogt*, Spektrum (2001), P.58-60
Moderne Betriebssysteme, *Andrew S. Tanenbaum*, Pearson (2009), P.874-877

Structure of a UNIX Process in Memory (5/6)

- With every function call, a data structure with the following contents is placed onto the stack:
 - Call parameters
 - Return address
 - Pointer to the calling function in the stack
- The functions also add (*push*) their local variables onto the stack
- When returning from from a function, the data structure of the function is removed (*pop*) from the stack




Sources

UNIX-Systemprogrammierung, *Helmut Herold*, Addison-Wesley (1996), P.345-347
Betriebssysteme, *Carsten Vogt*, Spektrum (2001), P.58-60
Moderne Betriebssysteme, *Andrew S. Tanenbaum*, Pearson (2009), P.874-877

Structure of a UNIX Process in Memory (6/6)

- The command `size` returns the size (in Bytes) of the text segment, data segment and BSS of program files
 - The contents of the text segment and data segment are included in the program files
 - All contents in the BSS are set to value 0 at process creation

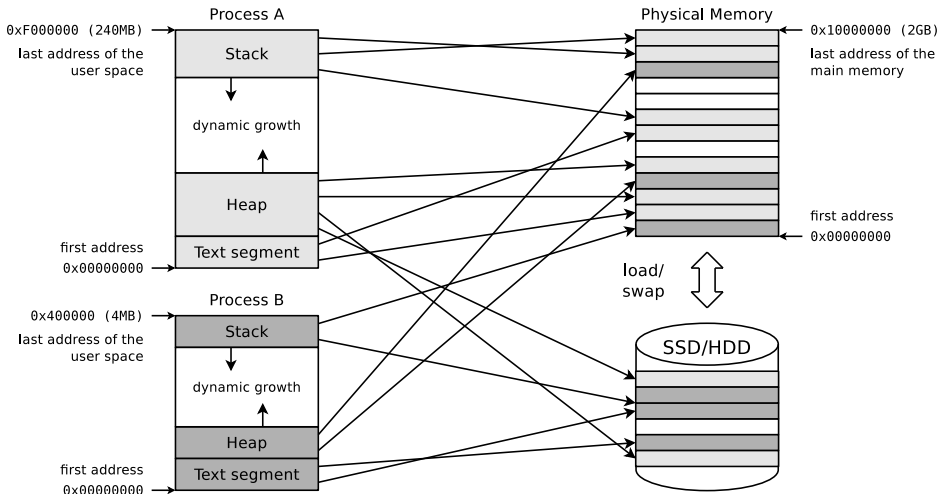
| \$ | size | /bin/c* | | | | |
|----|--------|---------|------|--------|-------|------------|
| | text | data | bss | dec | hex | filename |
| | 46480 | 620 | 1480 | 48580 | bdc4 | /bin/cat |
| | 7619 | 420 | 32 | 8071 | 1f87 | /bin/chacl |
| | 55211 | 592 | 464 | 56267 | dbcb | /bin/chgrp |
| | 51614 | 568 | 464 | 52646 | cda6 | /bin/chmod |
| | 57349 | 600 | 464 | 58413 | e42d | /bin/chown |
| | 120319 | 868 | 2696 | 123883 | 1e3eb | /bin/cp |
| | 131911 | 2672 | 1736 | 136319 | 2147f | /bin/cpio |



Sources

UNIX-Systemprogrammierung, *Helmut Herold*, Addison-Wesley (1996), P.345-347
Betriebssysteme, *Carsten Vogt*, Spektrum (2001), P.58-60
Moderne Betriebssysteme, *Andrew S. Tanenbaum*, Pearson (2009), P.874-877

Remember: Virtual Memory (Slide Set 5)




Source: http://cseweb.ucsd.edu/classes/wi11/cse141/Slides/19_VirtualMemory.key.pdf

Processes are stored in physical memory (here 2 GB) by virtual memory, not in a continuous manner and not always in main memory

User Mode and Kernel Mode

- x86-compatible CPUs implement 4 privilege levels
 - Objective: Improve stability and security
 - Each process is assigned to a ring permanently and can not free itself from this ring



Implementation of the privilege levels

- The register CPL (Current Privilege Level) stores the current privilege level
- Source: Intel 80386 Programmer's Reference Manual 1986
<http://css.csail.mit.edu/6.858/2012/readings/i386.pdf>
- In ring 0 (= **kernel mode**) runs the kernel
 - Here, processes have full access to the hardware
 - The kernel can also address physical memory (\implies Real Mode)
- In ring 3 (= **user mode**) run the applications
 - Here, processes can only access virtual memory (\implies Protected Mode)

Modern operating systems use only 2 privilege levels (rings)

Reason: Some hardware architectures (e.g. Alpha, PowerPC, MIPS) implement only 2 levels

System Calls (1/2)

We already know...

All processes outside the operating system kernel are only allowed to access their own virtual memory


- If a user-mode process must carry out a higher privileged task (e.g. access hardware), it can tell this the kernel via a **system call**
 - A system call is a function call in the operating system that triggers a switch from user mode to kernel mode (\implies **context switch**)

Context switch

- A process passes the control over the CPU to the kernel and is suspended until the request is completely processed
 - After the system call, the kernel returns the control over the CPU to the user-mode process
 - The process continues its execution at the point, where the context switch was previously requested
-
- The functionality of a system call is provided in the kernel
 - Thus, outside of the address space of the calling process

System Calls (2/2)

- **System calls** are the interface, which provides the operating system to the user mode processes
 - System calls enable the user mode programs among others to create and manage processes and files and to access the hardware




Simply stated...

A system call is a request from a user mode process to the kernel in order to use a service of the kernel

Comparison between System Calls and Interrupts

Interrupts (\implies slide set 3) are triggered by events outside user-mode processes



Example of a System Call: `ioctl()`

- This way, Linux programs call device-specific instructions
 - `ioctl()` enables processes to communicate with and control of:
 - Character devices (Mouse, keyboard, printer, terminals, ...)
 - Block devices (SSD/HDD, CD/DVD drive, ...)
- Syntax:

```
ioctl (File descriptor, request code number, integer value or pointer to data);
```

- Typical application scenarios of `ioctl()`:
 - Format floppy track
 - Initialize modem or sound card
 - Eject CD
 - Retrieve status and link information of the WLAN interface
 - Access sensors via the Inter-Integrated Circuit (I²C) data bus

Helpful overviews about system calls

Linux: <http://www.digilife.be/quickreferences/qrc/linux%20system%20call%20quick%20reference.pdf>

Linux: <http://syscalls.kernelgrok.com>

Linux: http://www.tutorialspoint.com/unix_system_calls/

Windows: <http://j00ru.vexillum.org/ntapi>

System Calls and Libraries

- Working directly with system calls is **unsafe** and the **portability is poor**
- Modern operating systems provide a library, which is logically located between the user mode processes and the kernel

Examples of such libraries

C Standard Library (UNIX), GNU C library glibc (Linux), C Library Implementationen (BSD), Native API `ntdll.dll` (Windows)

- The library is responsible for:
 - Handling the communication between user mode processes and kernel
 - Context switching between user mode and kernel mode
- Advantages which result in using a library:
 - Increased **portability**, because there is no or very little need for the user mode processes to communicate directly with the kernel
 - Increased **security**, because the user mode processes can not trigger the context switch to kernel mode for themselves





Image Source: Wikipedia
(Shmuel Csaba Otto Traian, CC-BY-SA-3.0)

Step by Step (1/4) – read(fd, buffer, nbytes);

- In step **1-3** stores the user mode process the parameters on the stack
- In **4** calls the user mode process the **library wrapper function** for read (\Rightarrow read nbytes from the file fd and store it inside buffer)
- In **5** stores the library wrapper function the system call number in the *accumulator register* EAX (32 bit) or RAX (64 bit)
 - The library wrapper function stores the parameters of the system call in the registers EBX, ECX and EDX (or for 64 bit: RBX, RCX and RDX)




Source of this example

Moderne Betriebssysteme, Andrew S. Tanenbaum, 3rd edition, Pearson (2009), P.84-89

Step by Step (2/4) – read(fd, buffer, nbytes);

- In **6**, the software interrupt (exception) 0x80 (decimal: 128) is triggered to switch from user mode to kernel mode
 - The software interrupt interrupts the program execution in user mode and enforces the execution of an exception handler in kernel mode




The kernel maintains the *System Call Table*, a list of all system calls

In this list, each system call is assigned to a unique number and an internal kernel function


Step by Step (3/4) – read(fd, buffer, nbytes);

- The called exception handler is a function in the kernel, which reads out the content of the EAX register
- The exception handler function calls in **7**, the corresponding kernel function from the system call table with the arguments, which are stored in the registers EBX, ECX and EDX
- In **8**, the system call is executed



Step by Step (4/4) – read(fd, buffer, nbytes);

- In **9**, the exception handler returns control back to the library, which triggered the software interrupt
- Next, this function returns in **10** back to the user mode process, in the way a normal function would have done it
- To complete the system call, the user mode process must clean up the stack in **11** just like after every function call
- The user process can now continue to operate



Example of a System Call in Linux

- System calls are called like library wrapper functions
 - The mechanism is similar for all operating systems
 - In a C program, no difference is visible

```
1 #include <syscall.h>
2 #include <unistd.h>
3 #include <stdio.h>
4 #include <sys/types.h>
5
6 int main(void) {
7     unsigned int ID1, ID2;
8
9     // System call
10    ID1 = syscall(SYS_getpid);
11    printf ("Result of the system call: %d\n", ID1);
12
13    // Wrapper function of the glibc, which calls the system call
14    ID2 = getpid();
15    printf ("Result of the wrapper function: %d\n", ID2);
16
17    return(0);
18 }
```

```
$ gcc SysCallBeispiel.c -o SysCallBeispiel
$ ./SysCallBeispiel
Result of the system call: 3452
Result of the wrapper function: 3452
```

Selection of System Calls

Process management

| | |
|---------|---|
| fork | Create a new child process |
| waitpid | Wait for the termination of a child process |
| execve | Replace a process by another one. The PID is kept |
| exit | Terminate a process |

File management

| | |
|-------|--|
| open | Open file for reading/writing |
| close | Close an open file |
| read | Read data from a file into the buffer |
| write | Write data from the buffer into a file |
| lseek | Reposition read/write file offset |
| stat | Determine the status of a file |

Directory management

| | |
|--------|---|
| mkdir | Create a new directory |
| rmdir | Remove an empty directory |
| link | Create a directory entry (link) to a file |
| unlink | Erase a directory entry |
| mount | Attach a file system to the file system hierarchy |
| umount | Detach a file system |

Miscellaneous

| | |
|-------|---|
| chdir | Change current directory |
| chmod | Change file permissions of a file |
| kill | Send signal to a process |
| time | Seconds since January 1st, 1970 („UNIX time“) |

Linux System Calls

- The list with the names of the system calls in the Linux kernel...
 - is located in the source code of kernel 2.6.x in the file:
arch/x86/kernel/syscall_table_32.S
 - is located in the source code of kernel 3.x, 4.x and 5.x in these files:
arch/x86/syscalls/syscall_[64|32].tbl or
arch/x86/entry/syscalls/syscall_[64|32].tbl

arch/x86/syscalls/syscall_32.tbl

```
...  
1      i386    exit      sys_exit  
2      i386    fork      sys_fork  
3      i386    read      sys_read  
4      i386    write     sys_write  
5      i386    open      sys_open  
6      i386    close     sys_close  
...
```

Tutorials how to implement own system calls

<https://www.kernel.org/doc/html/v4.14/process/adding-syscalls.html>
<https://brennan.io/2016/11/14/kernel-dev-ep3/>
<https://medium.com/@jeremyphilemon/adding-a-quick-system-call-to-the-linux-kernel-cad55b421a7b>
<https://medium.com/@ssreehari/implementing-a-system-call-in-linux-kernel-4-7-1-6f98250a8c38>
<http://tldp.org/HOWTO/Implement-Sys-Call-Linux-2.6-i386/index.html>
<http://www.ibm.com/developerworks/library/l-system-calls/>