

8th Slide Set

Operating Systems

Prof. Dr. Christian Baun


Frankfurt University of Applied Sciences
(1971–2014: Fachhochschule Frankfurt am Main)
Faculty of Computer Science and Engineering
christianbaun@fb2.fra-uas.de

Learning Objectives of this Slide Set

- At the end of this slide set You know/understand. . .
 - what steps the **dispatcher** carries out for switching between processes
 - what **scheduling** is
 - how **preemptive scheduling** and **non-preemptive scheduling** works
 - the functioning of several common **scheduling methods**
 - how **scheduling in modern operating systems** works in detail

In SS2019 I erased all scheduling algorithms (SJF/SRTF/LJF/LRTF/HRRN) from my course material that require to know how long it takes for each process until its termination. In other words, these algorithms need to now, how long is the execution time of each process. In practice this is almost never the case (\Rightarrow **unrealistic**)

Exercise sheet 8 repeats the contents of this slide set which are relevant for these learning objectives



Process Switching – The Dispatcher (1/2)

- Tasks of multitasking operating systems are among others:
 - **Dispatching**: Switching of the CPU during a process switch
 - **Scheduling**: Determination of the point in time when the process switch occurs and of the execution order of the processes
- The **dispatcher** carries out the state transitions of the processes

We already know...

- During process switching, the dispatcher removes the CPU from the `running` process and assigns it to the process, which is the first one in the queue
- For transitions between the states `ready` and `blocked`, the dispatcher removes the corresponding process control blocks from the status lists and accordingly inserts them new
- Transitions from or to the state `running` always imply a switch of the process, which is currently executed by the CPU

If a process switches into the state `running` or from the state `running` to another state, the dispatcher needs to...

- back up the context (register contents) of the executed process in the process control block
- assign the CPU to another process
- import the context (register contents) of the process, which will be executed next, from its process control block

Process Switching – The Dispatcher (2/2)

The system idle process

- Windows operating systems since Windows NT ensure that the CPU is assigned to a process at any time
- If no process is in the state ready, the **system idle process** gets the CPU assigned
- The system idle process is always active and has the lowest priority
- Due to the system idle process, the scheduler must never consider the case that no active process exists
- Since Windows 2000, the system idle process puts the CPU into a power-saving mode
- For each CPU core (in hyperthreading systems for each logical CPU), exists a system idle process

Image Name	User Name	CPU	Mem Usage
System Idle Process	SYSTEM	99	16 K
smssv.exe	SYSTEM	00	4,236 K
wscntfy.exe	BNC	00	1,904 K
svchost.exe	LOCAL SERVICE	00	4,292 K
taskmgr.exe	BNC	00	3,816 K
svchost.exe	NETWORK SERVICE	00	3,320 K
explorer.exe	BNC	00	12,876 K
wuauclt.exe	SYSTEM	00	8,196 K
svchost.exe	SYSTEM	00	25,212 K
alg.exe	LOCAL SERVICE	00	3,348 K
svchost.exe	NETWORK SERVICE	00	3,960 K
svchost.exe	SYSTEM	00	4,604 K
lsass.exe	SYSTEM	00	4,220 K
services.exe	SYSTEM	00	3,056 K
winlogon.exe	SYSTEM	00	1,352 K
csrss.exe	SYSTEM	00	2,872 K
wmiprvse.exe	SYSTEM	00	4,988 K
smss.exe	SYSTEM	00	356 K
msiexec.exe	SYSTEM	00	5,504 K

☐ Show processes from all users End Process

Processes: 20 CPU Usage: 0% Commit Charge: 97M / 3943M

<https://unix.stackexchange.com/questions/361245/what-does-an-idle-cpu-process-do>

„In Linux, one idle task is created for every CPU and locked to that processor; whenever there's no other process to run on that CPU, the idle task is scheduled. Time spent in the idle tasks appears as "idle" time in tools such as top. . .“

Scheduling Criteria and Scheduling Strategies

- During scheduling, the operating system specifies the execution order of the processes in the state ready
- **No scheduling strategy. . .**
 - **is optimally suited for each system**
 - **can take all scheduling criteria optimal into account**
 - Scheduling criteria are among others CPU load, response time (latency), turnaround time, throughput, efficiency, real-time behavior (compliance with deadlines), waiting time, overhead, fairness, consideration of priorities, even resource utilization. . .
- When choosing a scheduling strategy, a **compromise** between the scheduling criteria must always be found

Non-preemptive and preemptive Scheduling

- 2 classes of scheduling strategies exist
 - **Non-preemptive scheduling** or **cooperative scheduling**
 - A process, which gets the CPU assigned by the scheduler, remains control over the CPU until its execution is finished or it gives the control back on a voluntary basis
 - Problematic: A process may occupy the CPU for as long as it wants

Examples: Windows 3.x, MacOS 8/9, Windows 95/98/Me (for 16-Bit processes)

- **Preemptive scheduling**
 - The CPU may be removed from a process before its execution is completed
 - If the CPU is removed from a process, it is paused until the scheduler again assigns the CPU to it
 - Drawback: Higher overhead compared with non-preemptive scheduling
 - The benefits of preemptive scheduling, especially the consideration of process priorities, outweighs the drawbacks

Examples: Linux, MacOS X, Windows 95/98/Me (for 32-Bit processes), Windows NT (incl. XP/Visa/7/8/10/11), FreeBSD

Impact on the overall Performance of a Computer

- This example demonstrates the impact of the scheduling method used on the overall performance of a computer
 - The processes P_A and P_B are to be executed one after the other
 - If a short-running process runs before a long-running process, the runtime and wanting time of the long process get **slightly worse**
 - If a long-running process runs before a short-running process, the runtime and wanting time of the short process get **significantly worse**

Process	CPU runtime
A	24 ms
B	2 ms

Execution order	Runtime		Average runtime	Waiting time		Average waiting time
	A	B		A	B	
P_A, P_B	24 ms	26 ms	$\frac{24+26}{2} = 25$ ms	0 ms	24 ms	$\frac{0+24}{2} = 12$ ms
P_B, P_A	26 ms	2 ms	$\frac{2+26}{2} = 14$ ms	2 ms	0 ms	$\frac{0+2}{2} = 1$ ms

Scheduling Methods

- Several scheduling methods (algorithms) exist
 - Each method tries to comply with the well-known scheduling criteria and principles in varying degrees
- Some scheduling methods:
 - **Priority-driven scheduling**
 - **First Come First Served (FCFS) = First In First Out (FIFO)**
 - ~~Last Come First Served (LCFS)~~
 - **Round Robin (RR)** with time quantum
 - ~~Shortest/Longest Job First (SJF/LJF)~~
 - ~~Shortest/Longest Remaining Time First (SRTF/LRTF)~~
 - ~~Highest Response Ratio Next (HRRN)~~
 - **Earliest Deadline First (EDF)**
 - **Fair-share scheduling**
 - ~~Static multilevel scheduling~~
 - **Multilevel feedback scheduling**
 - **Completely Fair Scheduler (CFS)**

Modern operating systems often implement several scheduling methods

- In Linux e.g. each processes is assigned to a specific scheduling method
- For „**real-time**“ processes...
 - SCHED_FIFO (priority-driven scheduling, non-preemptive)
 - SCHED_RR (preemptive)
 - SCHED_DEADLINE (EDF scheduling, preemptive)
- For „**normal**“ processes...
 - SCHED_OTHER (default Linux time-sharing scheduling) implemented as...
 - Multilevel Feedback Scheduling (until Kernel 2.4)
 - O(1) scheduler (Kernel 2.6.0 until 2.6.22)
 - Completely Fair Scheduler (since Kernel 2.6.23)

```
$ ps a | grep okular
359675 pts/2    Sl          0:04 okular  bts_WS2122_slideset_08_en.pdf
$ chrt -p 361443
pid 361443's current scheduling policy: SCHED_OTHER
pid 361443's current scheduling priority: 0
```

```
SCHED_OTHER:      chrt -o -p PRIO PID
SCHED_FIFO:       chrt -f -p PRIO PID
SCHED_RR:         chrt -r -p PRIO PID
SCHED_DEADLINE:   chrt -d -sched-runtime NS -sched-deadline NS -sched-period NS 0 PID
```

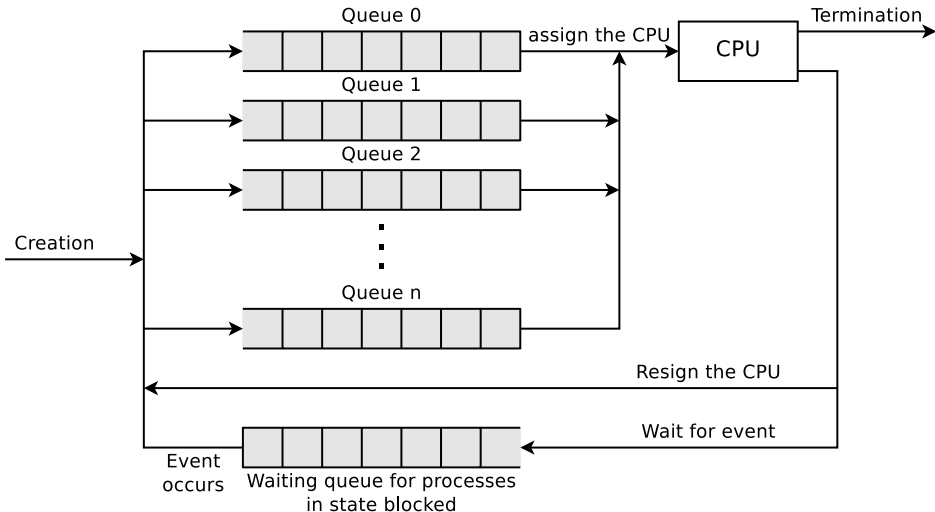
„A SCHED_DEADLINE task should receive *runtime* microseconds of execution time every *period* microseconds, and these *runtime* microseconds are available within *deadline* microseconds from the beginning of the period.“

Source: <https://www.kernel.org/doc/Documentation/scheduler/sched-deadline.txt>

Priority-driven Scheduling

- Processes are executed according to their priority (= importance or urgency)
- The highest priority process in state ready gets the CPU assigned
 - The priority may depend on various criteria, such as static (assigned) priority level, required resources, rank of the user, demanded real-time criteria,...
- Can be **preemptive** and **non-preemptive**
- The priority values can be assigned **static** or **dynamic**
 - Static priorities remain unchanged throughout the lifetime of a process, and are often used in real-time systems
 - Dynamic priorities are adjusted from time to time
⇒ **Multilevel feedback scheduling** (see slide 21)
- Risk of (static) priority-driven scheduling: Processes with low priority values may starve (⇒ **this is not fair**)
- Priority-driven scheduling can be used for interactive systems


Priority-driven Scheduling



Priority-driven Scheduling – Example

- 4 processes shall be processed on a single CPU/core system
- All processes are at time point 0 in state ready
- Execution order of the processes as Gantt chart (timeline)

Process	CPU time	Priority
A	8 ms	3
B	4 ms	15
C	7 ms	8
D	13 ms	4



- The CPU time is the time that the process needs to access the CPU to complete its execution
- Runtime = „lifetime“ = time period between the creation and the termination of a process = (CPU time + waiting time)

Runtime of the processes

Process	A	B	C	D
Runtime	32	4	11	24

$$\text{Avg. runtime} = \frac{32+4+11+24}{4} = 17.75 \text{ ms}$$

Waiting time of the processes

Process	A	B	C	D
Waiting time	24	0	4	11

$$\text{Avg. waiting time} = \frac{24+0+4+11}{4} = 9.75 \text{ ms}$$


First Come First Served (FCFS)

- Works according to the principle **First In First Out** (FIFO)
- This scheduling method is similar to a waiting line of customers in a store
- Running processes are not interrupted
 - It is **non-preemptive scheduling**
- FCFS is **fair**
 - All processes are executed
- The **average waiting time may be very high** under certain circumstances
 - Processes with short execution time may need to wait for a long time if processes with long execution times have arrived before
- FCFS/FIFO can be used for batch processing (\implies slide set 1)
- FIFO is used in Linux for non-preemptive „real-time“ processes

First Come First Served – Example

- 4 processes shall be processed on a single CPU/core system
- Execution order of the processes as Gantt chart (timeline)

Process	CPU time	Creation time
A	8 ms	0 ms
B	4 ms	1 ms
C	7 ms	3 ms
D	13 ms	5 ms



- The CPU time is the time that the process needs to access the CPU to complete its execution
- Runtime = „lifetime“ = time period between the creation and the termination of a process = (CPU time + waiting time)

Runtime of the processes

Process	A	B	C	D
Runtime	8	11	16	27

$$\text{Avg. runtime} = \frac{8+11+16+27}{4} = 15.5 \text{ ms}$$


Waiting time of the processes

Process	A	B	C	D
Waiting time	0	7	9	14

$$\text{Avg. waiting time} = \frac{0+7+9+14}{4} = 7.5 \text{ ms}$$

Round Robin – RR (1/2)

- Time slices with a fixed duration are specified
- The processes are queued in a cyclic queue according to the FIFO principle
 - The first process of the queue gets the CPU assigned for the duration of a time slice
 - After the expiration of the time slice, the process gets the CPU resigned and it is positioned at the end of the queue
 - Whenever a process is completed successfully, it is removed from the queue
 - New processes are inserted at the end of the queue
- The CPU time is distributed **fair** among the processes
- RR with time slice size ∞ behaves like FCFS




Round Robin – RR (2/2)

- The longer the execution time of a process is, the more rounds are required for its complete execution
- The size of the time slices influences the performance of the system
 - The shorter they are, the more process switches must take place
⇒ Increased overhead
 - The longer they are, the more gets the simultaneousness lost
⇒ The system hangs/becomes *jerky*
- The size of the time slices is usually in single or double-digit millisecond range
- **Prefers processes with short execution time**
- **Preemptive scheduling method**
- Round Robin scheduling can be used for interactive systems
- Round Robin is used in Linux for preemptive „real-time“ processes

Round Robin – Example

- 4 processes shall be processed on a single CPU/core system
- All processes are at time point 0 in state ready
- Time quantum $q = 1$ ms
- Execution order of the processes as Gantt chart (timeline)

Process	CPU time
A	8 ms
B	4 ms
C	7 ms
D	13 ms



- The CPU time is the time that the process needs to access the CPU to complete its execution
- Runtime = „lifetime“ = time period between the creation and the termination of a process = (CPU time + waiting time)

Runtime of the processes

Process	A	B	C	D
Runtime	26	14	24	32

$$\text{Avg. runtime} = \frac{26+14+24+32}{4} = 24 \text{ ms}$$

Waiting time of the processes

Process	A	B	C	D
Waiting time	18	10	17	19

$$\text{Avg. waiting time} = \frac{18+10+17+19}{4} = 16 \text{ ms}$$


Earliest Deadline First (EDF)

- Objective: processes should comply with their (deadlines) when possible
- Processes in ready state are **arranged according to their deadline**
 - The process with the closest deadline gets the CPU assigned next
- The queue is reviewed and reorganized whenever. . .
 - a new process switches into state ready
 - or an active process terminates
- Can be implemented as **preemptive and non-preemptive scheduling**
 - Preemptive EDF can be used in real-time operating systems
 - Non-preemptive EDF can be used for batch processing
- EDF is used in Linux for preemptive „real-time“ processes

Earliest Deadline First – Example

- 4 processes shall be processed on a single CPU/core system
- All processes are at time point 0 in state ready
- Execution order of the processes as Gantt chart (timeline)

Process	CPU time	Deadline
A	8 ms	25
B	4 ms	18
C	7 ms	9
D	13 ms	34



- The CPU time is the time that the process needs to access the CPU to complete its execution
- Runtime = „lifetime“ = time period between the creation and the termination of a process = (CPU time + waiting time)

Runtime of the processes

Process	A	B	C	D
Runtime	19	11	7	32


$$\text{Avg. runtime} = \frac{19+11+7+32}{4} = 17.25 \text{ ms}$$

Waiting time of the processes

Process	A	B	C	D
Waiting time	11	7	0	19

$$\text{Avg. waiting time} = \frac{11+7+0+19}{4} = 9.25 \text{ ms}$$

Fair-Share



- **Fair-Share** distributes the available resources between groups of processes in a fair manner
- Special feature:
 - The computing time is allocated to the users and not the processes
 - The computing time, which is allocated to a user, is independent from the number of his processes
- Users get resource shares

Fair share is often used in cluster and grid systems


Fair share is implemented in job schedulers and meta-schedulers (e.g. SUN/Oracle/Univa/Altair Grid Engine) for assigning the jobs to resources in grid sites distributing jobs between grid sites

Multilevel Feedback Scheduling (1/2)

- It is **impossible to predict the execution time precisely in advance**
 - Solution: Processes, which utilized much execution time in the past, get **sanctioned**
- **Multilevel feedback scheduling** works with multiple queues
 - Each queue has a different priority or time multiplex (e.g. 70%:15%:10%:5%)
- Each new process is added to the top queue
 - This way it has the highest priority
- Each queue uses Round Robin
 - If a process returns the CPU on voluntary basis, it is added to the same queue again
 - If a process utilized its entire time slice, it is inserted in the next lower queue, with has a lower priority
 - The priorities are therefore **dynamically** assigned with this method
- Multilevel feedback scheduling is **preemptive Scheduling**

Multilevel Feedback Scheduling (2/2)

- Benefit:
 - **No complicated estimations!**
 - New processes are quickly assigned to a priority category
- **Prefers new processes** over older (longer-running) processes
- Processes with many Input and output operations are preferred because they are inserted in the original queue again when they resigns the CPU on voluntary basis \Rightarrow This way they keep their priority value
- Older, longer-running processes are delayed




Source: William Stallings. Operating Systems. 4th edition. Prentice Hall (2001). P.413

Many modern operating systems use variants of multilevel feedback scheduling for the scheduling of the processes.
Examples: Linux for „normal“ processes (until Kernel 2.4), Mac OS X, FreeBSD, NetBSD and the Windows NT family

Completely Fair Scheduler (Linux since 2.6.23) – Part 1/4


- The kernel implements a CFS scheduler for every CPU core and maintains a variable **vruntime** (virtual runtime) for every SCHED_OTHER process
 - The value represents a virtual processor runtime in nanoseconds



- vruntime indicates how long the particular process has already used the CPU core
 - The process with the lowest vruntime gets access to the CPU core next
- The management of the processes is done using a **red-black tree** (self-balancing binary search tree)
 - The processes are sorted in the tree by their vruntime values

Completely Fair Scheduler (Linux since 2.6.23) – Part 2/4

- Aim: All processes should get a similar (fair) share of computing time of the CPU core they are assigned to
 \Rightarrow For n processes, each process should get $1/n$ of the CPU time




- If a process got the CPU core assigned, it can run until its vruntime value has reached the targeted portion of $1/n$ of the available CPU time
- The scheduler aims for an equal vruntime value for all processes

The CFS scheduler only takes care of the scheduling of the „normal“ (non-real-time) processes that are assigned to the scheduling method `SCHED_OTHER`

Completely Fair Scheduler (Linux since 2.6.23) – Part 3/4

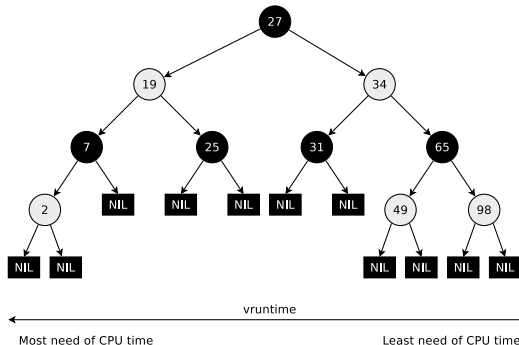
- The values are the keys of the inner nodes
- leaf nodes (NIL nodes) have no keys and contain no data
- NIL stands for *none*, *nothing*, *null*, which means it is a null value or null pointer



- For fairness reasons, the scheduler assigns the CPU core next to the leftmost process in the tree
- If a process gets replaced from the CPU core, the vruntime value is increased by the time the process did run on the CPU core

Completely Fair Scheduler (Linux since 2.6.23) – Part 4/4

- The nodes (processes) in the tree move continuously from right to left
⇒ fair distribution of CPU resources



- The scheduler takes into account the static process priorities (`nice` values) of the processes
- The `vruntime` values are weighted differently depending on the `nice` value
 - In other words: The virtual clock can run at different speeds

Classic and modern Scheduling Methods

	Scheduling NP	P	Fair	CPU time must be known	Takes priorities into account
Priority-driven scheduling	X	X	no	no	yes
First Come First Served = FIFO	X		yes	no	no
Last Come First Served	X	X	no	no	no
Round Robin		X	yes	no	no
Shortest/Longest Job First	X		no	yes	no
Shortest Remaining Time First		X	no	yes	no
Longest Remaining Time First		X	no	yes	no
Highest Response Ratio Next	X		yes	yes	no
Earliest Deadline First	X	X	yes	no	no
Fair-share		X	yes	no	no
Static multilevel scheduling		X	no	no	yes (static)
Multilevel feedback scheduling		X	yes	no	yes (dynamic)
O(1)-Scheduler		X	yes	no	yes
Completely Fair Scheduler		X	yes	no	yes

- NP = non-preemptive scheduling, P = preemptive scheduling
- A scheduling method is „fair“ when each process gets the CPU assigned at some point
- It is impossible to calculate the execution time precisely in advance

Linux 2.6.0 until 2.6.22 implements the **O(1) scheduler**. It does not play a role here for time reasons
<https://www.ibm.com/developerworks/library/l-scheduler/index.html>