

Exercise Sheet 6

Exercise 1 (File Systems)

1. Describe the information inodes store.
2. Name three examples of metadata in the file system.
3. Describe what a cluster in the file system is.
4. Describe how a UNIX file system (e.g. ext2/3), which does not implement extents, can address more than 12 clusters.
5. Describe how directories in the Linux file systems are technically implemented.
6. Explain why moving a large file within a file system is always faster than copying it.
7. Name one advantage and one drawback of small clusters in the file system compared with large clusters.
8. Do DOS/Windows file systems differentiate between uppercase and lowercase?
☐ Yes ☐ No
9. Do UNIX file systems differentiate between uppercase and lowercase?
☐ Yes ☐ No
10. Do modern operating systems accelerate requests to stored data with a cache in the main memory.
☐ Yes ☐ No
11. Most operating systems operate according to the principle...
☐ write-back ☐ write-through
12. Name one advantage and one disadvantage of a cache in the main memory, which is used by the operating system to accelerate the requests to stored data.
13. Explain what an absolute path name is.
14. Explain what a relative path name is.
15. `/var/log/messages` is an/a...
☐ absolute path name ☐ relative path name

16. `BTS_Vorlesung/Vorlesung_05/bts_slides_05_en.tex` is an/a...
☐ absolute path name ☐ relative path name
17. `Documents/MasterThesis/thesis.tex` is an/a...
☐ absolute path name ☐ relative path name
18. `/home/<username>/Mail/inbox/` is an/a...
☐ absolute path name ☐ relative path name
19. Describe what information the boot sector (also called boot block) of a file system stores.
20. Describe what information the super block of a file system stores.
21. Explain why some file systems (e.g. ext2/3) do combine the clusters of the file system to block groups.
22. Describe what the File Allocation Table (FAT) is and the information it stores.
23. Describe the objective of the journal in a journaling file system.
24. Describe a benefit of using a journaling file system compared with using a file system without a journal.
25. Name the three values that are required to store an extent.
26. Describe the benefit of using extents compared with direct addressing of the clusters.
27. Name one Linux file system that implements extents.
28. Name one Windows file system that implements extents.
29. When a file within a copy-on-write (COW) file system gets modified, the old clusters in the file system that need to be modified are...
☐ are overwritten with the new modifications.
☐ are preserved (not changed).
☐ are copied into new clusters, where the modifications are made.
☐ are erased, by removing the cluster address in the inode.
- (Hint: Two answers are correct!)
30. Describe the result of defragmenting a file system.
31. Describe the sort of data processing that is maximum accelerated by defragmenting.
32. Describe the scenario where defragmenting is useful.

Exercise 2 (File Systems)

Please mark for each statement about file systems, whether the statement is true or false.

Statement	true	false
Inodes store all metadata of files.		
File systems address clusters and not blocks of the storage medium or storage drive.		
The smaller the clusters are, the more overhead for large files occur.		
The bigger the clusters are, the lesser capacity is lost due to internal fragmentation.		
In UNIX, file extensions have always been of great significance.		
Modern file systems operate so much efficient that buffering by the operating system is no longer common.		
Absolute path names describe the complete path from the root to the file.		
The separator in path names is identical for all operating systems.		
An advantage of block groups is that the inodes are physically located close to the clusters, they address.		
For each cluster in the file system, an entry exists in the FAT.		
Because of the Master File Table in NTFS, fragmentation cannot occur.		
The journal of journaling file systems reduces the number of write operations.		
Journaling file systems narrow down the data, which need to be checked during the consistency check.		
When using journaling file systems, a loss of data is impossible.		
If metadata and file contents are journaled both, all write operations are carried out twice.		
Extents cause lesser overhead compared with block addressing.		

Exercise 3 (Pattern Comparison and Data Analysis)

1. Name (or describe) one useful application for the command `sed`.
2. Create a file `sedtest.txt` with the following content:

Line 1
Line 2
Line 3
Line 4
Line 5

Line 6

Insert with `sed` 3 blanks at the beginning of each line.

3. Print out with `sed` the lines 2 to 5 of the file `sedtest.txt`.
4. Remove with `sed` each second line of the file `sedtest.txt`.
5. Create a file `htmlcode.txt` with the following content:

```
<a href="BTSWS2019/index.html">Operating Systems (OpSys)</a><p>
<b>This is a <i>HTML file</i></b><br>
<h2>This is a headline<h2>
```

Remove with `sed` all HTML tags from the file `htmlcode.html`.

6. Create a file `umlaute.txt` with the following content:

Bäume, Äpfel, Bücher, Übertreibung
Töpfe, Öffentlichkeit, Straße, Spaß

Modify with `sed` all umlauts in the file `umlaute.txt` into „ae“, „oe“, „ue“, „Ae“, „Oe“, „Ue“ and „ss“.

7. Create a file `bundesliga_08_0405.txt` with the results of the 8th match day of the season 2004/2005:

Schalke	- Bochum	3 : 2	61500 spectators
Bielefeld	- Stuttgart	0 : 2	22700 spectators
Dortmund	- Nürnberg	2 : 2	73500 spectators
Leverkusen	- Hamburg	3 : 0	22500 spectators
Freiburg	- Mainz	1 : 2	24000 spectators
Kaiserslautern	- Berlin	0 : 2	30500 spectators
Wolfsburg	- Mönchengladbach	2 : 1	26500 spectators
Rostock	- Hannover	1 : 3	16500 spectators
Bremen	- München	1 : 2	42000 spectators

8. Name (or describe) one useful application for the command `awk`.
9. Determine with `awk` all matches, which had more than 30000 spectators.
10. Determine with `awk` all matches, which had less than 50000 spectators and where the home team won.
11. Determine with `awk` for each game the sum of the scored goals.
12. Determine with `awk` in which city the most spectators visited the match and print out the result this way:

The most spectators were in CITY (NUMBER).